

WELCOME!

In this project, you must apply EDA processes for the development of predictive models. Handling outliers, domain knowledge and feature engineering will be challenges.

Also, this project aims to improve your ability to implement algorithms for Multi-Class Classification. Thus, you will have the opportunity to implement many algorithms commonly used for Multi-Class Classification problems.

Before diving into the project, please take a look at the determines and tasks.

Determines

The 2012 US Army Anthropometric Survey (ANSUR II) was executed by the Natick Soldier Research, Development and Engineering Center (NSRDEC) from October 2010 to April 2012 and is comprised of personnel representing the total US Army force to include the US Army Active Duty, Reserves, and National Guard. In addition to the anthropometric and demographic data described below, the ANSUR II database also consists of 3D whole body, foot, and head scans of Soldier participants. These 3D data are not publicly available out of respect for the privacy of ANSUR II participants. The data from this survey are used for a wide range of equipment design, sizing, and tariffing applications within the military and has many potential commercial, industrial, and academic applications.

The ANSUR II working databases contain 93 anthropometric measurements which were directly measured, and 15 demographic/administrative variables explained below. The ANSUR II Male working database contains a total sample of 4,082 subjects. The ANSUR II Female working database contains a total sample of 1,986 subjects.

DATA DICT: <https://data.world/datamil/ansur-ii-data-dictionary/workspace/file?filename=ANSUR+II+Databases+Overview.pdf> (<https://data.world/datamil/ansur-ii-data-dictionary/workspace/file?filename=ANSUR+II+Databases+Overview.pdf>)

To achieve high prediction success, you must understand the data well and develop different approaches that can affect the dependent variable.

Firstly, try to understand the dataset column by column using pandas module. Do research within the scope of domain (body scales, and race characteristics) knowledge on the internet to get to know the data set in the fastest way.

You will implement **Logistic Regression**, **Support Vector Machine**, **XGBoost**, **Random Forest** algorithms. Also, evaluate the success of your models with appropriate performance metrics.

At the end of the project, choose the most successful model

Tasks

1. Exploratory Data Analysis (EDA)

- Import Libraries, Load Dataset, Exploring Data

i. Import Libraries

**ii. Ingest Data **

iii. Explore Data

iv. Outlier Detection

v. Drop unnecessary features

2. Data Preprocessing

- Scale (if needed)
- Separate the data frame for evaluation purposes

3. Multi-class Classification

- Import libraries
- Implement SVM Classifier
- Implement Decision Tree Classifier
- Implement Random Forest Classifier
- Implement XGBoost Classifier
- Compare The Models

Import Libraries

Besides Numpy and Pandas, you need to import the necessary modules for data visualization, data preprocessing, Model building and tuning.

In [1]:

```
#!pip install pyforest  
#!pip install cufflinks
```

In [2]:

```

import pyforest
import plotly
import cufflinks as cf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

#Enabling the offline mode for interactive plotting locally
from plotly.offline import download_plotlyjs,init_notebook_mode,plot,iplot
init_notebook_mode(connected=True)
cf.go_offline()

#To display the plots
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score, cross_validate
from sklearn.metrics import accuracy_score, f1_score, recall_score, precision_score
from sklearn.metrics import make_scorer
from sklearn.metrics import classification_report, confusion_matrix, plot_confusion_matrix
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import roc_curve, auc
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

pd.set_option('display.max_rows', 1000)
pd.set_option('display.max_columns', 1000)
pd.set_option('display.width', 1000)

```

Ingest Data from links below and make a dataframe

- Soldiers Male : <https://query.data.world/s/h3pbhckz5ck4rc7qmt2wlknlnn7esr>
(<https://query.data.world/s/h3pbhckz5ck4rc7qmt2wlknlnn7esr>)
- Soldiers Female : <https://query.data.world/s/sq27zz4hawg32yfxksqwijxmpwmyng>
(<https://query.data.world/s/sq27zz4hawg32yfxksqwijxmpwmyng>)

In [3]:

```

df_m = pd.read_csv('https://query.data.world/s/h3pbhckz5ck4rc7qmt2wlknlnn7esr', encoding='latin1')
df_f = pd.read_csv('https://query.data.world/s/sq27zz4hawg32yfxksqwijxmpwmyng')

print("df_m shape" , df_m.shape)
print("df_f shape" , df_f.shape)

```

```

df_m shape (4082, 108)
df_f shape (1986, 108)

```

In [4]:

```
df_m.head()
```

Out[4]:

	subjectid	abdominalextensiondepthsitting	acromialheight	acromionradialelength	anklecircumf
0	10027	266	1467	337	
1	10032	233	1395	326	
2	10033	287	1430	341	
3	10092	234	1347	310	
4	10093	250	1585	372	

In [5]:

```
df_f.head()
```

Out[5]:

	SubjectId	abdominalextensiondepthsitting	acromialheight	acromionradialelength	anklecircumf
0	10037	231	1282	301	
1	10038	194	1379	320	
2	10042	183	1369	329	
3	10043	261	1356	306	
4	10051	309	1303	308	

EDA

- Drop unnecessary columns
- Drop DODRace class if value count below 500 (we assume that our data model can't learn if it is below 500)

In [6]:

```
df_columns_names=pd.concat([pd.DataFrame(df_f.columns), pd.DataFrame(df_m.columns)], axis =  
df_columns_names
```

Out[6]:

	0	0
0	SubjectId	subjectid
1	abdominalextensiondepthsitting	abdominalextensiondepthsitting
2	acromialheight	acromialheight
3	acromionradialelength	acromionradialelength
4	anklecircumference	anklecircumference
5	axillaheight	axillaheight
6	ballofootcircumference	ballofootcircumference
7	ballofootlength	ballofootlength
8	biacromialbreadth	biacromialbreadth
9	bicepscircumferenceflexed	bicepscircumferenceflexed


In [7]:

```
df = pd.concat([df_m, df_f])
df
```

Out[7]:

	subjectid	abdominalextensiondepthsitting	acromialheight	acromionradialelength	anklecirc
0	10027.0		266	1467	337
1	10032.0		233	1395	326
2	10033.0		287	1430	341
3	10092.0		234	1347	310
4	10093.0		250	1585	372
...
1981	NaN		285	1392	335
1982	NaN		262	1324	301
1983	NaN		260	1334	318
1984	NaN		205	1293	302
1985	NaN		238	1346	308

6068 rows × 109 columns



In [8]:

```
df.columns
```

Out[8]:

```
Index(['subjectid', 'abdominalextensiondepthsitting', 'acromialheight', 'acromionradialelength', 'anklecircumference', 'axillaheight', 'balloffootcircumference', 'balloffootlength', 'biacromialbreadth', 'bicepscircumferenceflexed',
      ...
      'PrimaryMOS', 'SubjectsBirthLocation', 'SubjectNumericRace', 'Ethnicity', 'DODRace', 'Age', 'Heightin', 'Weightlbs', 'WritingPreference', 'SubjectId'], dtype='object', length=109)
```

In [9]:

```
df[["subjectid", "SubjectId"]]
```

Out[9]:

	subjectid	SubjectId
0	10027.0	NaN
1	10032.0	NaN
2	10033.0	NaN
3	10092.0	NaN
4	10093.0	NaN
...
1981	NaN	29501.0
1982	NaN	29502.0
1983	NaN	29503.0
1984	NaN	29511.0
1985	NaN	920103.0

6068 rows × 2 columns

In [10]:

```
df.drop(["subjectid", "SubjectId"], axis=1, inplace = True)
```

In [11]:

```
df.shape
```

Out[11]:

```
(6068, 107)
```

In [12]:

```
df.isnull().sum().any(), df.duplicated().sum()
```

Out[12]:

(True, 0)

In [13]:

```
df.isnull().sum()
```

Out[13]:

abdominalextensiondepthsitting	0
acromialheight	0
acromionradialelength	0
anklecircumference	0
axillaheight	0
balloffootcircumference	0
balloffootlength	0
biacromialbreadth	0
bicepscircumferenceflexed	0
bicristalbreadth	0
bideltoidbreadth	0
bimalleolarbreadth	0
bitracionchinarc	0
bitrationsubmandibulararc	0
bizygomaticbreadth	0
buttockcircumference	0
buttockdepth	0
buttockheight	0

In [14]:

```
NaN_list = []

for columns in df.columns:
    if df[columns].isnull().sum() > 0:
        print("{name} = {quantity}".format(name = columns, quantity = df[columns].isnull().sum()))
        NaN_list.append(columns)
```

Ethnicity = 4647

In [15]:

```
df.drop(["Ethnicity"], axis = 1, inplace = True)
```

In [16]:

```
df.isnull().sum().any()
```

Out[16]:

False

In [17]:

```
df.shape
```

Out[17]:

```
(6068, 106)
```

In [18]:

```
df[["SubjectNumericRace", "DODRace"]]
```

Out[18]:

	SubjectNumericRace	DODRace
0	1	1
1	1	1
2	2	2
3	1	1
4	2	2
...
1981	3	3
1982	3	3
1983	2	2
1984	3	3
1985	3	3

6068 rows × 2 columns

In [19]:

```
df[df["SubjectNumericRace"] == df["DODRace"].any().any()]
```

Out[19]:

```
True
```

In [20]:

```
df.drop(["SubjectNumericRace"], axis = 1, inplace = True)
```

In [21]:

```
df.shape
```

Out[21]:

```
(6068, 105)
```

In [22]:

```
df.DODRace.value_counts()
```

Out[22]:

```
1    3792
2    1298
3     679
4     188
6      59
5      49
8       3
Name: DODRace, dtype: int64
```

In [23]:

```
df = df[df.DODRace < 4]

# df = df[df["DODRace"].isin([1,2,3])]
```

In [24]:

```
df.DODRace.value_counts()
```

Out[24]:

```
1    3792
2    1298
3     679
Name: DODRace, dtype: int64
```

In [25]:

```
df.shape
```

Out[25]:

```
(5769, 105)
```

In [26]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 5769 entries, 0 to 1985
Columns: 105 entries, abdominalextensiondepthsitting to WritingPreference
dtypes: int64(97), object(8)
memory usage: 4.7+ MB
```

In [27]:

```

#!pip install researchpy
import researchpy as rp

object_columns = df.select_dtypes(include='object')
rp.summary_cat(object_columns)

```

Out[27]:

	Variable	Outcome	Count	Percent
0	Gender	Male	3899	67.59
1		Female	1870	32.41
2	Date	5-Apr-12	43	0.75
3		5-Mar-12	43	0.75
4		28-Feb-12	42	0.73
5		26-May-11	40	0.69
6		21-Feb-12	40	0.69
7		6-Mar-12	40	0.69
8		26-Mar-12	39	0.68
9		15-Mar-11	39	0.68

In [28]:

```

for columns in object_columns:
    print(f"{columns:<25}:", df[columns].nunique())           # <25 ile
# print(f"{columns} has {df[columns].nunique()} unique value") # farklı b

```

```

Gender           : 2
Date             : 253
Installation     : 12
Component        : 3
Branch           : 3
PrimaryMOS       : 281
SubjectsBirthLocation : 136
WritingPreference : 3

```

In [29]:

```
df.drop(["Date", "Installation", "Component", "Branch", "PrimaryMOS", "WritingPreference"],
```

In [30]:

```
df.shape
```

Out[30]:

```
(5769, 99)
```

In [31]:

```
df.Gender.value_counts()
```

Out[31]:

```
Male      3899  
Female    1870  
Name: Gender, dtype: int64
```

In [32]:

```
for columns in df.select_dtypes(exclude = [object]):  
    print(columns)
```

abdominalextensiondepthsitting
acromialheight
acromionradialelength
anklecircumference
axillaheight
balloffootcircumference
balloffootlength
biacromialbreadth
bicepscircumferenceflexed
bicristalbreadth
bideltoidebreadth
bimalleolarbreadth
bitragionchinarc
bitragionsubmandibulararc
bizygomaticbreadth
buttockcircumference
buttockdepth
buttockheight
buttockkneelength
buttockpopliteallength
calfcircumference
cervicaleheight
chestbreadth
chestcircumference
chestdepth
chestheight
crotchheight
crotchlenthomphalion
crotchlenthposterioromphalion
earbreadth
earlength
earprotrusion
elbowrestheight
eyeheightsitting
footbreadthhorizontal
footlength
forearmcenterofgriplength
forearmcircumferenceflexed
forearmforearmbreadth
forearmhandlength
functionalleglength
handbreadth
handcircumference
handlength
headbreadth
headcircumference
headlength
heelanklecircumference
heelbreadth
hipbreadth
hipbreadthsitting
iliocristaleheight
interpupillarybreadth
interscyei
interscyeii
kneeheightmidpatella

kneeheightsitting
lateral femoral epicondyle height
lateral malleolus height
lower thigh circumference
menton sellion length
neck circumference
neck circumference base
overhead fingertip reach sitting
palm length
popliteal height
radial styloid length
shoulder circumference
shoulder elbow length
shoulder length
sitting height
sleeve length spinewrist
sleeve outseam
span
stature
suprasternal height
tenth rib height
thigh circumference
thigh clearance
thumb tip reach
tibial height
tragion top of head
trochanterion height
vertical trunk circumference cusa
waist back length
waist breadth
waist circumference
waist depth
waist front length sitting
waist height to navel
weight kg
wrist circumference
wrist height
DOD Race
Age
Height in
Weight lbs

In [33]:

```
df[["weightkg", "Weightlbs"]]
```

Out[33]:

	weightkg	Weightlbs
0	815	180
1	726	160
2	929	205
3	794	175
4	946	213
...
1981	832	180
1982	717	150
1983	762	168
1984	632	133
1985	610	132

5769 rows × 2 columns

In [34]:

```
df[["Heightin", "stature"]]
```

Out[34]:

	Heightin	stature
0	71	1776
1	68	1702
2	68	1735
3	66	1655
4	77	1914
...
1981	67	1687
1982	63	1613
1983	66	1644
1984	63	1616
1985	66	1641

5769 rows × 2 columns

In [35]:

```
df.drop(['Weightlbs', 'Heightin'], axis = 1, inplace = True)
```

In [36]:

```
df["weightkg"].head()
```

Out[36]:

```
0    815
1    726
2    929
3    794
4    946
Name: weightkg, dtype: int64
```

In [37]:

```
df["weightkg"] = df["weightkg"]/10
```

In [38]:

```
df["weightkg"].head()
```

Out[38]:

```
0    81.5
1    72.6
2    92.9
3    79.4
4    94.6
Name: weightkg, dtype: float64
```

In [39]:

```
df.shape
```

Out[39]:

```
(5769, 97)
```


In [40]:

```

# function for set text color of positive
# values in Dataframes

def color_red(val):
    """
    Takes a scalar and returns a string with
    the css property ``color: red`` for positive
    strings, black otherwise.
    """
    if (1 > val > 0.9) or (-0.9 > val > -1):
        color = 'red'

    elif val==1:
        color='blue'

    else:
        color = 'black'

    return 'color: %s' % color

pd.DataFrame(df).corr().style.applymap(color_red)

```

Out[40]:

	abdominalextensiondepthsitting	acromialheight	acromionradialelength	anklec
abdominalextensiondepthsitting	1.000000	0.351934	0.312919	
acromialheight	0.351934	1.000000	0.868267	
acromionradialelength	0.312919	0.868267	1.000000	
anklecircumference	0.518896	0.504673	0.416051	
axillaheight	0.280824	0.987115	0.857391	
ballofootcircumference	0.456729	0.693952	0.604208	
ballofootlength	0.332593	0.797793	0.725966	
biacromialbreadth	0.417617	0.733288	0.667377	
bicepscircumferenceflexed	0.691126	0.522740	0.452499	

In [41]:

```

df_corr = df.corr()

count = 0
feature = []
collinear = []
for columns in df_corr.columns:

    for i in df_corr.index:

        if (1 > df_corr[columns][i] > 0.9) or (-0.9 > df_corr[columns][i] > -1) :
            feature.append(columns)
            collinear.append(i)
            count += 1
            #print(f"multicollinearity alert in between {columns} - {i}")    # bununla a

print("Number of strong corelated features:", count)

```

Number of strong corelated features: 286

In [42]:

```

df_corr = pd.DataFrame([feature, collinear], index=["feature", "collinear"]).T
df_corr

```

Out[42]:

	feature	collinear
0	abdominalextensiondepthsitting	waistcircumference
1	abdominalextensiondepthsitting	waistdepth
2	acromialheight	axillaheight
3	acromialheight	cervicaleheight
4	acromialheight	chestheight
5	acromialheight	iliocristaleheight
6	acromialheight	kneeheightsitting
7	acromialheight	stature
8	acromialheight	suprasternaleheight
9	acromialheight	tenthribheight

In [43]:

```
df.head()
```

Out[43]:

	abdominalextensiondepthsitting	acromialheight	acromionradialelength	anklecircumference	ax
0	266	1467	337	222	
1	233	1395	326	220	
2	287	1430	341	230	
3	234	1347	310	230	
4	250	1585	372	247	

In [44]:

```
df2 = df.copy()
```

In [45]:

```
df = pd.get_dummies(df, drop_first =True)
```

In [46]:

```
df.shape
```

Out[46]:

```
(5769, 231)
```

In [47]:

```
df.isnull().any().any()
```

Out[47]:

```
False
```

In [48]:

```

#!/pip install movecolumn
import movecolumn as mc

df = mc.MoveToLast(df, 'DODRace')

```

In [49]:

```
df["DODRace"] = df.DODRace.map({1 : 0, 2 : 1, 3 : 2})
```

In [50]:

```
df.head()
```

Out[50]:

	abdominalextensiondepthsitting	acromialheight	acromionradialelength	anklecircumference	ax
0	266	1467	337	222	
1	233	1395	326	220	
2	287	1430	341	230	
3	234	1347	310	230	
4	250	1585	372	247	

DATA Preprocessing

- In this step we divide our data to X(Features) and y(Target) then ,
- To train and evaluation purposes we create train and test sets,
- Lastly, scale our data if features not in same scale. Why?

In [51]:

```
X = df.drop("DODRace", axis = 1)
y = df.DODRace
```

In [52]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify = y ,rand
```

Modelling

- Fit the model with train dataset
- Get predict from vanilla model on both train and test sets to examine if there is over/underfitting
- Apply GridsearchCV for both hyperparameter tuning and sanity test of our model.
- Use hyperparameters that you find from gridsearch and make final prediction and evaluate the result according to chosen metric.

In [53]:

```
def eval_metric(model, X_train, y_train, X_test, y_test):  
    y_train_pred = model.predict(X_train)  
    y_pred = model.predict(X_test)  
  
    print("Test_Set")  
    print(confusion_matrix(y_test, y_pred))  
    print(classification_report(y_test, y_pred))  
  
    print()  
  
    print("Train_Set")  
    print(confusion_matrix(y_train, y_train_pred))  
    print(classification_report(y_train, y_train_pred))
```

1. Logistic model

Vanilla Logistic Model

In [58]:

```

from sklearn.pipeline import Pipeline

operations = [("scaler", MinMaxScaler()), ("log_pipeline", LogisticRegression(class_weight=
pipe_model=Pipeline(steps=operations)

pipe_model.fit(X_train, y_train)

y_pred = pipe_model.predict(X_test)

eval_metric(pipe_model, X_train, y_train, X_test, y_test)

```

Test_Set

```

[[626  25 107]
 [  6 241  13]
 [ 27   9 100]]

```

	precision	recall	f1-score	support
0	0.95	0.83	0.88	758
1	0.88	0.93	0.90	260
2	0.45	0.74	0.56	136
accuracy			0.84	1154
macro avg	0.76	0.83	0.78	1154
weighted avg	0.87	0.84	0.85	1154

Train_Set

```

[[2638  60 336]
 [  30 968  40]
 [  56  21 466]]

```

	precision	recall	f1-score	support
0	0.97	0.87	0.92	3034
1	0.92	0.93	0.93	1038
2	0.55	0.86	0.67	543
accuracy			0.88	4615
macro avg	0.81	0.89	0.84	4615
weighted avg	0.91	0.88	0.89	4615

In [59]:

```
f1_2 = make_scorer(f1_score, average=None, labels=[2])
precision_2 = make_scorer(precision_score, average=None, labels=[2])
recall_2 = make_scorer(recall_score, average=None, labels=[2])

scores = cross_validate(pipe_model, X_train, y_train, scoring = {'f1_2':f1_2,
                                                                'precision_2':precision_2,
                                                                'recall_2':recall_2},
                                                                cv = 10)

df_scores = pd.DataFrame(scores, index = range(1, 11))
df_scores.mean()[2:]
```

Out[59]:

```
test_f1_2          0.621272
test_precision_2   0.514768
test_recall_2      0.784579
dtype: float64
```

https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html (https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)

Logistic Model GridsearchCV

In [150]:

```
operations = [("scaler", MinMaxScaler()), ("log_pipeline", LogisticRegression(class_weight=
pipe_model=Pipeline(steps=operations))

f1_2 = make_scorer(f1_score, average=None, labels=[2])

C = np.logspace(-1, 5, 20)
penalty = ["l1", "l2"]
class_weight= ["balanced", None]
solver = ["liblinear", "saga"]

param_grid = {'log_pipeline__C': C,
              'log_pipeline__penalty': penalty,
              'log_pipeline__class_weight':class_weight,
              'log_pipeline__solver': solver }

LOGpipe_model_grid = GridSearchCV(pipe_model, param_grid, scoring = f1_2, cv = 5, n_jobs =
```

In [151]:

```
LOGpipe_model_grid.fit(X_train, y_train)
```

Out[151]:

```
GridSearchCV(cv=5,
             estimator=Pipeline(steps=[('scaler', MinMaxScaler()),
                                       ('log_pipeline',
                                        LogisticRegression(class_weight='balanced',
                                                           max_iter=10000,
                                                           random_state=4
                                                           2))]),
             n_jobs=-1,
             param_grid={'log_pipeline__C': array([1.00000000e-01, 2.0691380
8e-01, 4.28133240e-01, 8.85866790e-01,
1.83298071e+00, 3.79269019e+00, 7.84759970e+00, 1.62377674e+01,
3.35981829e+01, 6.95192796e+01, 1.43844989e+02, 2.97635144e+02,
6.15848211e+02, 1.27427499e+03, 2.63665090e+03, 5.45559478e+03,
1.12883789e+04, 2.33572147e+04, 4.83293024e+04, 1.00000000e+05]),
                       'log_pipeline__class_weight': ['balanced', None],
                       'log_pipeline__penalty': ['l1', 'l2'],
                       'log_pipeline__solver': ['liblinear', 'saga']},
             scoring=make_scorer(f1_score, average=None, labels=[2]))
```

In [152]:

```
LOGpipe_model_grid.best_score_
```

Out[152]:

```
0.6861932935133976
```

In [153]:

```
LOGpipe_model_grid.best_params_
```

Out[153]:

```
{'log_pipeline__C': 7.847599703514611,
 'log_pipeline__class_weight': 'balanced',
 'log_pipeline__penalty': 'l2',
 'log_pipeline__solver': 'liblinear'}
```


In [154]:

```
eval_metric(LOGpipe_model_grid, X_train, y_train, X_test, y_test)
```

Test_Set

```
[[698 16 44]
 [ 14 244 2]
 [ 42 9 85]]
```

	precision	recall	f1-score	support
0	0.93	0.92	0.92	758
1	0.91	0.94	0.92	260
2	0.65	0.62	0.64	136
accuracy			0.89	1154
macro avg	0.83	0.83	0.83	1154
weighted avg	0.89	0.89	0.89	1154

Train_Set

```
[[2869 36 129]
 [ 38 984 16]
 [ 112 24 407]]
```

	precision	recall	f1-score	support
0	0.95	0.95	0.95	3034
1	0.94	0.95	0.95	1038
2	0.74	0.75	0.74	543
accuracy			0.92	4615
macro avg	0.88	0.88	0.88	4615
weighted avg	0.92	0.92	0.92	4615

In [219]:


```

from yellowbrick.classifier import ROCAUC

model = LOGpipe_model_grid
visualizer = ROCAUC(model)

visualizer.fit(X_train, y_train)      # Fit the training data to the visualizer
visualizer.score(X_test, y_test)     # Evaluate the model on the test data
visualizer.show();

```



2. SVC

In []:

```

from sklearn.svm import SVC

```

Vanilla SVC model

In [81]:

```
# from sklearn.pipeline import Pipeline

operations = [("scaler", MinMaxScaler()), ("SVC_pipeline", SVC(random_state=42))]
SVCpipe_model=Pipeline(steps=operations)

SVCpipe_model.fit(X_train, y_train)

y_pred = SVCpipe_model.predict(X_test)

eval_metric(SVCpipe_model, X_train, y_train, X_test, y_test)
```

```
Test_Set
[[743  13   2]
 [ 31 229   0]
 [ 84   9  43]]
```

	precision	recall	f1-score	support
1	0.87	0.98	0.92	758
2	0.91	0.88	0.90	260
3	0.96	0.32	0.48	136
accuracy			0.88	1154
macro avg	0.91	0.73	0.76	1154
weighted avg	0.89	0.88	0.86	1154

```
Train_Set
[[3011  17   6]
 [ 111 924   3]
 [ 350  29 164]]
```

	precision	recall	f1-score	support
1	0.87	0.99	0.93	3034
2	0.95	0.89	0.92	1038
3	0.95	0.30	0.46	543
accuracy			0.89	4615
macro avg	0.92	0.73	0.77	4615
weighted avg	0.90	0.89	0.87	4615

SVC Model GridsearchCV

In [137]:

```

operations = [("scaler", MinMaxScaler()), ("SVC_pipeline", SVC(probability=True, class_weight='balanced'))]
pipe_model=Pipeline(steps=operations)

f1_2 = make_scorer(f1_score, average=None, labels=[2])
#precision_2 = make_scorer(precision_score, average=None, labels=[2])
#recall_2 = make_scorer(recall_score, average=None, labels=[2])

#scoring = {'SVC_f1_2':f1_2, 'SVC_precision_2':precision_2, 'SVC_recall_2':recall_2}

param_grid = {'SVC_pipeline__C': [0.01, 0.05, 0.1, 0.3, 1],
              'SVC_pipeline__gamma': ["scale", "auto", 0.2, 0.3],
              'SVC_pipeline__kernel': ['rbf', 'linear']}

SVCpipe_model_grid = GridSearchCV(pipe_model, param_grid,
                                  scoring = f1_2, verbose=2, n_jobs = -1)

```

In [138]:

```
SVCpipe_model_grid.fit(X_train, y_train)
```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

Out[138]:

```

GridSearchCV(estimator=Pipeline(steps=[('scaler', MinMaxScaler()),
                                       ('SVC_pipeline',
                                        SVC(class_weight='balanced',
                                             probability=True,
                                             random_state=42))])),
             n_jobs=-1,
             param_grid={'SVC_pipeline__C': [0.01, 0.05, 0.1, 0.3, 1],
                         'SVC_pipeline__gamma': ['scale', 'auto', 0.2, 0.3],
                         'SVC_pipeline__kernel': ['rbf', 'linear']},
             scoring=make_scorer(f1_score, average=None, labels=[2]),
             verbose=2)

```

In [139]:

```
eval_metric(SVCpipe_model_grid, X_train, y_train, X_test, y_test)
```

```
Test_Set
[[636  19 103]
 [ 13 237  10]
 [ 31   8  97]]
      precision    recall  f1-score   support

    0         0.94        0.84        0.88         758
    1         0.90        0.91        0.90         260
    2         0.46        0.71        0.56         136

 accuracy         0.84         1154
 macro avg         0.76         1154
weighted avg         0.87         1154
```

```
Train_Set
[[2632   53 349]
 [  32 970   36]
 [  52  16 475]]
      precision    recall  f1-score   support

    0         0.97        0.87        0.92        3034
    1         0.93        0.93        0.93        1038
    2         0.55        0.87        0.68         543

 accuracy         0.88         4615
 macro avg         0.82         4615
weighted avg         0.91         4615
```

In [140]:

```
SVCpipe_model_grid.best_params_
```

Out[140]:

```
{'SVC_pipeline__C': 1,
 'SVC_pipeline__gamma': 'scale',
 'SVC_pipeline__kernel': 'linear'}
```

In [141]:

```
SVCpipe_model_grid.best_score_
```

Out[141]:

```
0.6131733023498092
```

In []:

3. RF

In [207]:

```
df2.head()
```

Out[207]:

	abdominalextensiondepthsitting	acromialheight	acromionradialelength	anklecircumference	ax
0	266	1467	337	222	
1	233	1395	326	220	
2	287	1430	341	230	
3	234	1347	310	230	
4	250	1585	372	247	

In [208]:

```
from sklearn.preprocessing import OrdinalEncoder

objectColumns = df2.select_dtypes("object").columns
objectColumns

encoder = OrdinalEncoder()
df2[objectColumns] = encoder.fit_transform(df2[objectColumns])
df2.head()
```

Out[208]:

	abdominalextensiondepthsitting	acromialheight	acromionradialelength	anklecircumference	ax
0	266	1467	337	222	
1	233	1395	326	220	
2	287	1430	341	230	
3	234	1347	310	230	
4	250	1585	372	247	

In [209]:

```
df2["DODRace"] = df2.DODRace.map({1 : 0, 2 : 1, 3 : 2})
```

In [210]:

```
X2 = df2.drop("DODRace", axis = 1)
y2 = df2.DODRace

X2_train, X2_test, y2_train, y2_test = train_test_split(X2, y2, test_size=0.2, stratify = y
```

Vanilla RF Model

In [211]:

```
rf_model = RandomForestClassifier(class_weight = "balanced", random_state=42)

rf_model.fit(X2_train, y2_train)

eval_metric(rf_model, X2_train, y2_train, X2_test, y2_test)
```

```
Test_Set
[[745  11   2]
 [ 79 181   0]
 [123   9   4]]
```

	precision	recall	f1-score	support
0	0.79	0.98	0.87	758
1	0.90	0.70	0.79	260
2	0.67	0.03	0.06	136
accuracy			0.81	1154
macro avg	0.78	0.57	0.57	1154
weighted avg	0.80	0.81	0.76	1154

```
Train_Set
[[3034   0   0]
 [   0 1038   0]
 [   0   0 543]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3034
1	1.00	1.00	1.00	1038
2	1.00	1.00	1.00	543
accuracy			1.00	4615
macro avg	1.00	1.00	1.00	4615
weighted avg	1.00	1.00	1.00	4615

RF Model GridsearchCV

In [212]:

```
model = RandomForestClassifier(class_weight = "balanced", random_state=42)

param_grid = {'n_estimators':[50, 64, 100, 128, 300, 400, 500],
              'max_features':[2, 3, 4, "auto"],
              'max_depth':[3, 6, 9, 12, 15],
              'min_samples_split':[5, 10, 15, 20]}

f1_2 = make_scorer(f1_score, average=None, labels=[2])

RF3_model_grid = GridSearchCV(model, param_grid, scoring = f1_2, n_jobs = -1, verbose = 2)
```

In [213]:

```
RF3_model_grid.fit(X2_train, y2_train)
```

Fitting 5 folds for each of 560 candidates, totalling 2800 fits

Out[213]:

```
GridSearchCV(estimator=RandomForestClassifier(class_weight='balanced',
                                              random_state=42),
             n_jobs=-1,
             param_grid={'max_depth': [3, 6, 9, 12, 15],
                         'max_features': [2, 3, 4, 'auto'],
                         'min_samples_split': [5, 10, 15, 20],
                         'n_estimators': [50, 64, 100, 128, 300, 400, 500]},
             scoring=make_scorer(f1_score, average=None, labels=[2]),
             verbose=2)
```

In [216]:

```
eval_metric(RF3_model_grid, X2_train, y2_train, X2_test, y2_test)
```

Test_Set

```
[[631  41  86]
 [ 36 209  15]
 [ 73  15 48]]
```

	precision	recall	f1-score	support
0	0.85	0.83	0.84	758
1	0.79	0.80	0.80	260
2	0.32	0.35	0.34	136
accuracy			0.77	1154
macro avg	0.65	0.66	0.66	1154
weighted avg	0.78	0.77	0.77	1154

Train_Set

```
[[2812  45 177]
 [ 38 990  10]
 [ 48  4 491]]
```

	precision	recall	f1-score	support
0	0.97	0.93	0.95	3034
1	0.95	0.95	0.95	1038
2	0.72	0.90	0.80	543
accuracy			0.93	4615
macro avg	0.88	0.93	0.90	4615
weighted avg	0.94	0.93	0.93	4615

In [217]:

```
RF3_model_grid.best_params_
```

Out[217]:

```
{'max_depth': 9,  
 'max_features': 'auto',  
 'min_samples_split': 15,  
 'n_estimators': 500}
```

In [218]:

```
RF3_model_grid.best_score_
```

Out[218]:

```
0.39336922857409373
```

4. XGBoost

In [60]:

```
from xgboost import XGBClassifier
```

Vanilla XGBoost Model

In [61]:

```
XGB_model = XGBClassifier(random_state=42)

XGB_model.fit(X_train, y_train)

eval_metric(XGB_model, X_train, y_train, X_test, y_test)
```

```
Test_Set
[[738  12   8]
 [ 28 231   1]
 [ 70  11 55]]
```

	precision	recall	f1-score	support
0	0.88	0.97	0.93	758
1	0.91	0.89	0.90	260
2	0.86	0.40	0.55	136
accuracy			0.89	1154
macro avg	0.88	0.76	0.79	1154
weighted avg	0.89	0.89	0.88	1154

```
Train_Set
[[3034  0   0]
 [  0 1038  0]
 [  0   0 543]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3034
1	1.00	1.00	1.00	1038
2	1.00	1.00	1.00	543
accuracy			1.00	4615
macro avg	1.00	1.00	1.00	4615
weighted avg	1.00	1.00	1.00	4615

XGBoost Model GridsearchCV

In [142]:

```
XGB_model = XGBClassifier(random_state=42)

param_grid = {"n_estimators": [50, 100, 200, 300],
              'max_depth': [2, 3, 4, 5],
              "learning_rate": [0.1, 0.2, 0.3],
              "subsample": [0.5, 0.8, 1],
              "colsample_bytree": [0.5, 0.7, 1]}

f1_2 = make_scorer(f1_score, average=None, labels=[2])

XGB_model_grid = GridSearchCV(XGB_model, param_grid, scoring = f1_2 , refit = True, verbose
```

In [143]:

```
XGB_model_grid.fit(X_train, y_train)
```

Fitting 5 folds for each of 432 candidates, totalling 2160 fits

Out[143]:

```
GridSearchCV(estimator=XGBClassifier(base_score=None, booster=None,
                                     callbacks=None, colsample_bylevel=None,
                                     colsample_bynode=None,
                                     colsample_bytree=None,
                                     early_stopping_rounds=None,
                                     enable_categorical=False, eval_metric=None,
                                     gamma=None, gpu_id=None, grow_policy=None,
                                     importance_type=None,
                                     interaction_constraints=None,
                                     learning_rate=None, max_bin=None,
                                     max_cat_to_one=None,
                                     max_depth=None, max_features=None,
                                     min_child_weight=None, missing=None,
                                     monotone_constraints=None,
                                     n_estimators=100, n_jobs=None,
                                     num_parallel_tree=None, predictor=None,
                                     random_state=42, reg_alpha=None,
                                     reg_lambda=None, ...),
             param_grid={'colsample_bytree': [0.5, 0.7, 1],
                         'learning_rate': [0.1, 0.2, 0.3],
                         'max_depth': [2, 3, 4, 5],
                         'n_estimators': [50, 100, 200, 300],
                         'subsample': [0.5, 0.8, 1]},
             scoring=make_scorer(f1_score, average=None, labels=[2]),
             verbose=2)
```

In [144]:

```
eval_metric(XGB_model_grid, X_train, y_train, X_test, y_test)
```

Test_Set

```
[[724 15 19]
 [ 21 235 4]
 [ 70 13 53]]
```

	precision	recall	f1-score	support
0	0.89	0.96	0.92	758
1	0.89	0.90	0.90	260
2	0.70	0.39	0.50	136
accuracy			0.88	1154
macro avg	0.83	0.75	0.77	1154
weighted avg	0.87	0.88	0.87	1154

Train_Set

```
[[3027 2 5]
 [ 11 1027 0]
 [ 45 5 493]]
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	3034
1	0.99	0.99	0.99	1038
2	0.99	0.91	0.95	543
accuracy			0.99	4615
macro avg	0.99	0.97	0.98	4615
weighted avg	0.99	0.99	0.99	4615

In [145]:

```
XGB_model_grid.best_params_
```

Out[145]:

```
{'colsample_bytree': 0.7,
 'learning_rate': 0.3,
 'max_depth': 2,
 'n_estimators': 300,
 'subsample': 0.5}
```

In [146]:

```
XGB_model_grid.best_score_
```

Out[146]:

```
0.5768772252628614
```

In []:

Final Model

- Choose the model that works best based on your chosen metric
- For final step, fit the best model with whole dataset to get better performance.

In [55]:

```
from sklearn.pipeline import Pipeline

operations = [("scaler", MinMaxScaler()), ("final_model", LogisticRegression(C = 7.84759970
                                                                              penalty = "l2
                                                                              max_iter=1000

final_model=Pipeline(steps=operations)

final_model.fit(X, y)
```

Out[55]:

```
Pipeline(steps=[('scaler', MinMaxScaler()),
                 ('final_model',
                  LogisticRegression(C=7.847599703514611,
                                     class_weight='balanced', max_iter=10000,
                                     random_state=42, solver='liblinear'))])
```

In [57]:

```
f1_2 = make_scorer(f1_score, average=None, labels=[2])
precision_2 = make_scorer(precision_score, average=None, labels=[2])
recall_2 = make_scorer(recall_score, average=None, labels=[2])

scores = cross_validate(final_model, X, y, scoring = {'f1_2':f1_2,
                                                    'precision_2':precision_2,
                                                    'recall_2':recall_2},
                      cv = 5)

df_scores = pd.DataFrame(scores, index = range(1, 6))
df_scores.mean()[2:]
```

Out[57]:

```
test_f1_2          0.632737
test_precision_2   0.620392
test_recall_2      0.656950
dtype: float64
```